

PARAG AGRAWAL
paraga@cs.stanford.edu
www.stanford.edu/~paraga

Home Address

110 McFarland Ct., #101
Stanford, CA 94305

Office Address

353 Serra Mall, #430
Stanford, CA 94305

EDUCATION

Stanford University

PhD Student, Computer Science Department
Master of Science in Computer Science
Advisor: Prof. Jennifer Widom

Sept. 2005 - Present
Expected graduation: Summer 2011
June 2007

Indian Institute of Technology, Bombay

Bachelor of Technology, Department of Computer Science and Engineering
CPI 9.48/10.00, Ranked 1st in class of 48

May 2005

RESEARCH EXPERIENCE

Uncertain Data Management with Provenance (Trio and Panda)

Focus on uncertain data integration.
Advisor: Prof. Jennifer Widom, Stanford University

Oct. 2005 - Present

Database Summarization

Scalable mining of high-dimensional data.
Mentors: Dr. Divesh Srivastava, Dr. Flip Korn and Dr. Lukasz Golab, At&t Labs Research

July 2010 - Sept. 2010

Data Cleaning (DMX Group)

Indexing string similarity functions.
Mentor: Dr. Raghav Kaushik, Microsoft Research

June 2009 - Sept. 2009

Scalable Data Management (PNUTS)

Weak consistency, secondary indexes, materialized views.
Mentors: Dr. Raghu Ramakrishnan and Dr. Brian Cooper, Yahoo! Research

June 2008 - Sept. 2008
June 2007 - Sept. 2007

Large Scale Data Analysis (Pig/Hadoop)

Stage scheduling for multi-query optimization.
Mentor: Dr. Chris Olston, Yahoo! Research

June 2007 - Sept. 2007

Blog Search

Algorithms for ranking blogs.
Mentor: Dr. Rakesh Agrawal, Microsoft Search Labs

July 2006 - Sept. 2006

TEACHING EXPERIENCE

CS346: Database System Implementation

Co-instructor for advanced graduate class with another PhD student.

Stanford University
Spring 2010

PUBLICATIONS

- RAMClouds: Scalable High-Performance Storage Entirely in DRAM** CACM 2011
John Ousterhout, Parag Agrawal, David Erickson, Christos Kozyrakis, Jacob Leverich, David Mazires, Subhasish Mitra, Aravind Narayanan, Guru Parulkar, Mendel Rosenblum, Stephen M. Rumble, Eric Stratmann, Ryan Stutsman Also in SIGOPS OSR 2009
- Generalized Uncertain Databases: First Steps** MUD Workshop 2010
Parag Agrawal, Jennifer Widom
- Foundations of Uncertain-Data Integration** VLDB 2010
Parag Agrawal, Anish Das Sarma, Jeffrey Ullman, Jennifer Widom
- On Indexing Error-Tolerant Set Containment** SIGMOD 2010
Parag Agrawal, Arvind Arasu, Raghav Kaushik
- Continuous Uncertainty in Trio** MUD Workshop 2009
Parag Agrawal, Jennifer Widom
- Trio-ER: The Trio System as a Workbench for Entity-Resolution** Tech. Rep. 2009
Parag Agrawal, Robert Ikeda, Hyunjung Park, Jennifer Widom
- Asynchronous View Maintenance for VLSD Databases** SIGMOD 2009
Parag Agrawal, Adam Silberstein, Brian Cooper, Utkarsh Srivastava, Raghu Ramakrishnan
- Confidence-Aware Join Algorithms** ICDE 2009
Parag Agrawal, Jennifer Widom
- Scheduling Shared Scans of Large Data Files** VLDB 2008
Parag Agrawal, Daniel Kifer, Christopher Olston
- Towards Special-Purpose Indexes and Statistics for Uncertain Data** MUD Workshop 2008
Anish Das Sarma, Parag Agrawal, Shubha Nabar, Jennifer Widom
- Trio-One: Layering Uncertainty and Lineage on a Conventional DBMS** CIDR 2007 (Demo)
Michi Mutsuzaki, Martin Theobald, Ander de Keijzer, Jennifer Widom, Parag Agrawal, Omar Benjelloun, Anish Das Sarma, Raghotham Murthy, Tomoe Sugihara
- Trio: A System for Data, Uncertainty, and Lineage** VLDB 2006 (Demo)
Parag Agrawal, Omar Benjelloun, Anish Das Sarma, Chris Hayworth, Shubha Nabar, Tomoe Sugihara, Jennifer Widom
-

ACADEMIC HONORS

- Winner of Yahoo! Key Scientific Challenges Program in 2009.
- Recipient of a 3-year Cisco Systems Stanford Graduate Fellowship.
- Awarded Institute Silver Medal for academic excellence at IIT Bombay.
- Represented IIT Bombay at ACM-ICPC World Finals in 2005.
- Gold Medalist at International Physics Olympiad (IPhO) in 2001. Ranked 4th in Theory.